

# Advanced Computer vision

August 15, 2024

## 0.1 ADVANCED COMPUTER VISION

```
[5]: !pip install torch torchvision
```

```
Requirement already satisfied: torch in c:\users\user\anaconda3\lib\site-  
packages (2.4.0)  
Requirement already satisfied: torchvision in c:\users\user\anaconda3\lib\site-  
packages (0.19.0)  
Requirement already satisfied: filelock in c:\users\user\anaconda3\lib\site-  
packages (from torch) (3.13.1)  
Requirement already satisfied: typing-extensions>=4.8.0 in  
c:\users\user\anaconda3\lib\site-packages (from torch) (4.11.0)  
Requirement already satisfied: sympy in c:\users\user\anaconda3\lib\site-  
packages (from torch) (1.12)  
Requirement already satisfied: networkx in c:\users\user\anaconda3\lib\site-  
packages (from torch) (3.2.1)  
Requirement already satisfied: jinja2 in c:\users\user\anaconda3\lib\site-  
packages (from torch) (3.1.4)  
Requirement already satisfied: fsspec in c:\users\user\anaconda3\lib\site-  
packages (from torch) (2024.3.1)  
Requirement already satisfied: setuptools in c:\users\user\anaconda3\lib\site-  
packages (from torch) (69.5.1)  
Requirement already satisfied: numpy in c:\users\user\anaconda3\lib\site-  
packages (from torchvision) (1.26.4)  
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in  
c:\users\user\anaconda3\lib\site-packages (from torchvision) (10.3.0)  
Requirement already satisfied: MarkupSafe>=2.0 in  
c:\users\user\anaconda3\lib\site-packages (from jinja2->torch) (2.1.3)  
Requirement already satisfied: mpmath>=0.19 in c:\users\user\anaconda3\lib\site-  
packages (from sympy->torch) (1.3.0)
```

```
[7]: import torch  
import torchvision  
import torchvision.transforms as transforms
```

```
[11]: transform = transforms.Compose(  
    [transforms.ToTensor(),  
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```

batch_size = 4

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)

testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

Files already downloaded and verified  
Files already downloaded and verified

```

[13]: import matplotlib.pyplot as plt
import numpy as np


```

```

[15]: def imshow(img):
    img = img/2 + 0.5
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

    detaiter = iter(trainloader)
    images, labels = next(detaiter)
    imshow(torchvision.utils.make_grid(images))
    print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))

```



```

frog ] cat  ] ship ] car  ]

```

```
[ ]:
```

```
[16]: import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) #flatten all dimensions
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

```
[ ]:
```

```
[19]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()

optimizer = optim.SGD(net.parameters(), lr=0.001, momentum = 0.9 )
```

```
[23]: for epoch in range(2):
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
```

```

outputs = net(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

# print statistics
running_loss+= loss.item()
if i % 2000 == 1999: # print every 2000 mini-batches
    print(f'[{epoch + 1}, {i + 1:5d}] loss: {running_loss / 2000:.3f}')
    running_loss = 0.0

print('Finished Training')

```

```

[1, 2000] loss: 1.244
[1, 4000] loss: 1.244
[1, 6000] loss: 1.241
[1, 8000] loss: 1.233
[1, 10000] loss: 1.238
[1, 12000] loss: 1.196
[2, 2000] loss: 1.155
[2, 4000] loss: 1.137
[2, 6000] loss: 1.160
[2, 8000] loss: 1.146
[2, 10000] loss: 1.135
[2, 12000] loss: 1.151
Finished Training

```


[ ]:

```

[25]: dataiter = iter(testloader)
images, labels = next(dataiter)

# print images
imshow(torchvision.utils.make_grid(images))
print('GroundTruth:', ' '.join(f'{classes[labels[j]]:5s}' for j in range(4)))

```



GroundTruth: frog plane car horse

**Ruth.O.Ajagunna**